

Dataset Preparation for Stable Diffusion Using Google Gemini



Stable Diffusion is a powerful generative artificial intelligence (AI) model that primarily uses a latent diffusion model (LDM) to produce high-quality, detailed images from text descriptions (text-to-image generation) and other image-based inputs.

- **Diffusion Process:** The core idea is based on diffusion models, which learn to iteratively reverse a process of gradually adding noise to an image.
 - **Forward Diffusion:** Gaussian noise is progressively added to the original image until it becomes pure random noise.
 - **Reverse Diffusion(Generation)** The model learns to reverse this process, starting from pure noise and using a U-Net noise predictor guided by a text prompt to "denoise" the image iteratively back into a coherent and detailed output image.
- **Latent Space:** Stable Diffusion operates in a lower-dimensional latent space (thanks to a Variational Autoencoder or VAE), rather than the high-dimensional pixel space. This significantly reduces computational requirements, making it more accessible.

Leveraging Google Gemini for Dataset Preparation

The key bottleneck in preparing a dataset is often the laborious and subjective process of manual image captioning. Google's Gemini model, being multimodal, is exceptionally useful for automating this step:

Gemini can process visual input (the image) and generate highly descriptive, natural-language text output (the caption), acting as a powerful and efficient auto-captioner.

How to Use Gemini for Captioning:

- Access the API: You can use the Gemini API (e.g., via the Google Gen AI SDK) which supports multimodal input.
 - How to Use Gemini for Captioning

Step 1: Get a Gemini API Key

To use Gemini programmatically, you need an API key:

Go to the Gemini developer console: <https://ai.google.dev>

Sign in with your Google account

Click Get API Key

- Input: You pass the image file (or a base64 encoded version of the image) along with a text prompt that instructs the model on what kind of caption to generate.

This Python code is a comprehensive script designed to extend a public dataset of Flat UI screenshots by generating Stable Diffusion-ready text captions for some of the images using the Google Gemini API, and then uploading the augmented dataset to Kaggle.

1. Install & import everything
2. Setup Kaggle access

3. Load a base dataset (Flat UI images)
4. Visualize sample images
5. Use Gemini to generate captions for some images
6. Combine old + new data into one dataset
7. Save it to disk in Kaggle format
8. Upload to Kaggle as a new dataset

Step 1: Installation and Imports

The initial steps ensure all necessary libraries are installed and imported.

- `datasets` : For loading and manipulating datasets, particularly from Hugging Face.
- `google-generativeai` : The official Google library for interacting with the Gemini API.
- `kaggle` : The command-line interface tool to interact with Kaggle for dataset upload.
- `pillow`, `matplotlib`, `pandas` : Standard libraries for image handling, plotting, and data manipulation.

```
In [ ]: # 1. Install everything
!pip install -q datasets google-generativeai kaggle pillow matplotlib pandas
```

```
In [ ]: # Import required libraries
from datasets import load_dataset, concatenate_datasets, Dataset, DatasetDict
from PIL import Image
import io, os, time, json
import google.generativeai as genai
import matplotlib.pyplot as plt
import subprocess
from IPython.display import clear_output
```

2. Kaggle API Setup

This block sets up the necessary authentication for the Kaggle CLI to upload the dataset.

- `from google.colab import files` : Imports the utility to handle file uploads in a Google Colab environment.
- `files.upload()` : Prompts the user to upload their kaggle.json file, which contains their API credentials.
- `!mkdir -p ~/.kaggle` : Creates a hidden directory named .kaggle in the home directory. This is the standard location where the Kaggle CLI expects to find the credentials file.
- `!cp kaggle.json ~/.kaggle/` : Copies the uploaded kaggle.json file into the newly created .kaggle directory.
- `!chmod 600 ~/.kaggle/kaggle.json` : Changes the file permissions to ensure only the owner (the user) can read and write to the file, which is a security requirement for the Kaggle CLI.

```
In [ ]: #2. Kaggle API setup
# You MUST upload your `kaggle.json` from https://www.kaggle.com/settings/account → API → Create Token
from google.colab import files

print("Upload your kaggle.json (download from https://www.kaggle.com/settings → API)")
files.upload() # Opens a file picker to upload kaggle.json

# Move kaggle.json to correct location and set correct permissions
!mkdir -p ~/.kaggle           # create Kaggle config folder if not exists
!cp kaggle.json ~/.kaggle/      # copy credentials there
!chmod 600 ~/.kaggle/kaggle.json # secure the file so only you can access it (required by Kaggle)

Upload your kaggle.json (download from https://www.kaggle.com/settings → API)
 Choose Files No file chosen          Upload widget is only available when the cell has been executed in the current browser session. Please
rerun this cell to enable.
Saving kaggle.json to kaggle (3).json
```

3. Load Original Dataset

This step loads the base dataset from the Hugging Face Hub.

```
old_dataset= load_dataset("bhomik7/flat-UI-dataset-small") : Loads the "flat-UI-dataset-small" dataset, which contains UI screenshots as binary image data. The old_dataset is typically a DatasetDict containing splits like 'train', 'test', etc.
```

```
In [ ]: # 3. Load original dataset
old_dataset = load_dataset("bhomik7/flat-UI-dataset-small")
#print(f"Original size: {len(old_dataset['train'])} images")
```

4. Show Samples

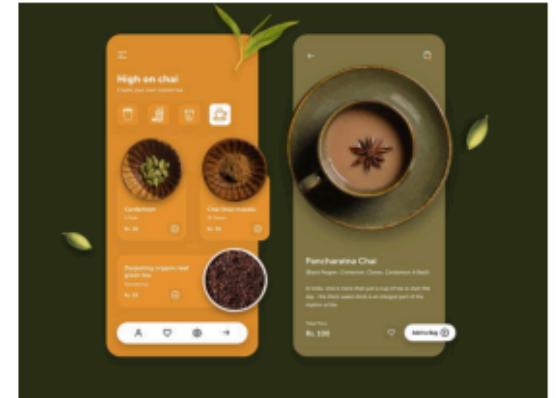
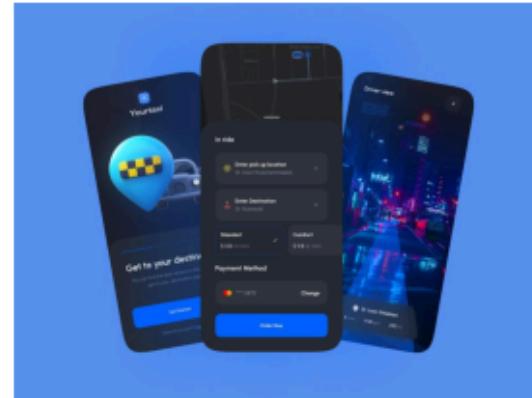
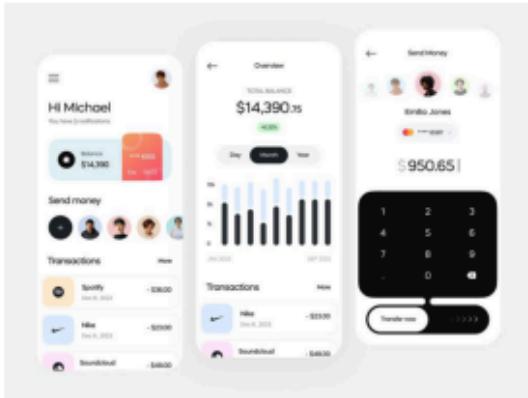
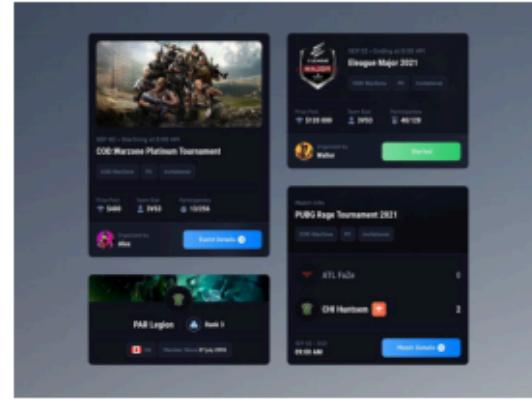
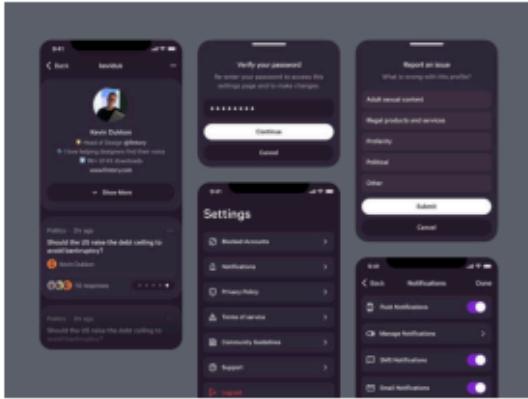
This defines a utility function to visually inspect a few images from the dataset.

- `def show_samples(ds, n=6): ...` : The function takes a dataset split (`ds`) and the number of samples (`n`) to show.
- It iterates, opens the image from its raw binary bytes using `io.BytesIO` and `PIL.Image.open()`, and plots them using `matplotlib.pyplot` in a grid format. This confirms the images loaded correctly.

```
In [ ]: # 4. Show samples
def show_samples(ds, n=6):
    """
    Display n sample images from the dataset.
    ds: Hugging Face dataset or dict with "image" keys (raw bytes).
    n: number of samples to show (default = 6)
    """
    plt.figure(figsize=(12,6))
    for i in range(n):
        # Convert raw image bytes into a displayable image
        img = Image.open(io.BytesIO(ds[i]["image"]))
        # Create subplot grid (2 rows x 3 columns)
        plt.subplot(2, 3, i+1)
        plt.imshow(img)
        plt.axis('off')
    plt.suptitle("Original Flat UI Samples")
    plt.show()
```

```
# Show example UI screenshots from the training set  
show_samples(old_dataset["train"])
```

Original Flat UI Samples



5. Gemini Setup and Captioning Function

This block configures the Gemini API access and defines the core function for generating captions.

- `genai.configure(api_key="...")` : Initializes the Gemini API client with the user's API key.

- `model = genai.GenerativeModel(...)` : Creates an instance of the model to be used:
 - "gemini-2.0-flash": A fast and capable multimodal model.
 - `system_instruction` : A powerful instruction to guide the model's behavior, directing it to output a short, precise description for a Stable Diffusion (SD) image generation model.
 - `generation_config` : Sets parameters like temperature (creativity) and `max_output_tokens`.
- `def gemini_caption(image_bytes) : ...`: This function takes raw image bytes and:
- Opens the image, converts it to RGB format and saves the image into an in-memory buffer (`io.BytesIO`) as a JPEG
- Uploads the buffered image data to the Gemini API using `genai.upload_file()`. This is necessary for the model to process the image.
- Calls `model.generate_content()` with both the uploaded image file and a specific prompt ("Give me a perfect SD prompt for this UI").
- `time.sleep(2)` : Pauses execution to prevent hitting API rate limits (sending too many requests too quickly).

Returns the generated text caption after stripping whitespace.

```
In [ ]: # 5. Configure Gemini API with your personal key (replace placeholder)
genai.configure(api_key="put_your_gemini_API_key_here")

# Define the Gemini model and its behavior
model = genai.GenerativeModel(
    "gemini-2.0-flash", # Gemini model name
    system_instruction=(
        "Describe this UI screenshot in <200 words so Stable Diffusion can recreate it exactly."
    ),
    generation_config={
        "temperature": 0.7,          # randomness in generation
        "max_output_tokens": 500     # maximum response length
    }
)

def gemini_caption(image_bytes):
    """
```

```

Send an image to Gemini and get back a high-quality caption
that can be used as a Stable Diffusion prompt.

"""
# Load the image from bytes and convert to RGB
img = Image.open(io.BytesIO(image_bytes)).convert("RGB")

# Re-encode it as JPEG (Gemini needs a file-like input)
buf = io.BytesIO()
img.save(buf, format="JPEG")
buf.seek(0)

# Upload to Gemini
file = genai.upload_file(buf, mime_type="image/jpeg")

# Ask Gemini to describe it for Stable Diffusion
response = model.generate_content([file, "Give me a perfect SD prompt for this UI."])

# Small delay to respect API rate limits
time.sleep(2)

# Return the clean text response
return response.text.strip()

```

6. Generate New Captions

This is the loop where the actual caption generation takes place.

- `for i in range(10):` : The script iterates through the first 10 images of the training split (the user notes this can be increased).
- Inside the loop
 - The raw image bytes are fetched.
 - `gemini_caption` is called to get a high-quality, Stable Diffusion-ready text description.
- The image bytes and the generated caption are stored in `new_images` and `new_texts` lists, respectively.
- A progress message is printed, showing the start of the caption.

```
In [ ]: new_images = []
new_texts = []

# Generate captions for first 10 images (can increase to 100 or 1000 later)
for i in range(10):
    print(f"Captioning image {i}...")

    # Get image bytes from dataset
    img_bytes = old_dataset["train"][i]["image"]

    # Use Gemini to create a caption (SD prompt)
    caption = gemini_caption(img_bytes)

    # Store the results
    new_images.append(img_bytes)
    new_texts.append(caption)

    # Show preview of the caption
    print(f"- {caption[:80]}...")
```

Captioning image 0...

- Here's a prompt to recreate the UI screenshot for Stable Diffusion:

**Detailed ...

Captioning image 1...

- A dark-themed UI design featuring tournament cards for various games. The top-le...

Captioning image 2...

- A close-up shot of a modern smartphone displaying a clean, minimalist UI for a "...

Captioning image 3...

- Here's a prompt to recreate the UI screens:

**UI Design, Mobile App Interface:*

Captioning image 4...

- A conceptual UI design for a ride-sharing app. Three phone screens are displayed...

Captioning image 5...

- UI design for a mobile application, split screen layout, with a dark olive green...

Captioning image 6...

- Here is a prompt that should recreate the image well:

A UI mockup featuring thr...

Captioning image 7...

- UI design showcasing four modal windows on a dark background. The top-left modal...

Captioning image 8...

- Here's a prompt to recreate the UI screenshot:

Prompt:

UI design of a car ...

Captioning image 9...

- Here's a Stable Diffusion prompt to recreate the UI screenshot, optimized for ac...

7. Build Final Dataset

The newly captioned data is merged with the original, uncaptioned data to create a single, unified dataset.

- `data = { ... }` : A dictionary is constructed to hold the data:
 - `"image"` : Contains the 10 captioned images plus all images from the original dataset.

- "text": Contains the 10 new captions plus an empty string ("") for every image from the original dataset (which had no pre-existing caption).
- `full_dataset = Dataset.from_dict(data)` : Creates a Hugging Face Dataset object from the combined dictionary.

```
In [ ]: #7. Build one unified dataset with images and captions
data = {
    # Combine new and old image bytes
    "image": new_images + [row["image"] for row in old_dataset["train"]],
    # Use new captions + empty strings for old (uncaptioned) images
    "text": new_texts + [""] * len(old_dataset["train"])
}

# Convert to a Hugging Face Dataset object
full_dataset = Dataset.from_dict(data)

print(f"FINAL SIZE: {len(full_dataset)} images")
```

FINAL SIZE: 20 images

8. Save to Folder for Kaggle

The `full_dataset` is saved into a local directory structure required for a Kaggle dataset upload.

- `SAVE_DIR = "flat_UI_extended_gemini"` : Defines the name of the directory.
- `os.makedirs(SAVE_DIR, exist_ok=True)` : Creates the output directory.
- The loop iterates through every item in the `full_dataset`:
 - Image Saving: It opens the image bytes and saves it as a JPEG file (`.jpg`) inside the `SAVE_DIR` . File names are padded with leading zeros (e.g., 00000.jpg).
 - Caption Saving: It saves the corresponding caption into a matching `.txt` file. This structure is common for image-text datasets.

```
In [ ]: # 8. Save the dataset locally (for Kaggle upload)
SAVE_DIR = "flat_UI_extended_gemini"
```

```

os.makedirs(SAVE_DIR, exist_ok=True)

for i in range(len(full_dataset)):
    img_bytes = full_dataset[i]["image"]
    caption = full_dataset[i]["text"] or "flat UI screenshot"

    # Save the image as a numbered JPG
    Image.open(io.BytesIO(img_bytes)).save(f"{SAVE_DIR}/{i:05d}.jpg")

    # Save the caption text (same filename, .txt extension)
    open(f"{SAVE_DIR}/{i:05d}.txt", "w").write(caption)

```

9. Kaggle Metadata

A `dataset-metadata.json` file is created. This file is crucial as the Kaggle CLI requires it to correctly title and license the uploaded dataset.

- The metadata includes a descriptive title, a unique ID (which includes the user's username, and the dataset name), and a license (here, the permissive CC0).

```

In [ ]: # 9. Create Kaggle metadata (dataset info file)
meta = {
    "title": "Flat UI + Gemini Captions (Stable Diffusion Ready)",
    "id": "bhomik7/flat-ui-dataset-extended",    # Kaggle dataset slug
    "licenses": [{"name": "cc0"}]                  # open license (public domain)
}

# Write metadata to JSON file (required by Kaggle)
json.dump(meta, open(f"{SAVE_DIR}/dataset-metadata.json", "w"), indent=2)

```

10. Kaggle Upload

This is the final, essential step where the local folder is published to Kaggle.

- `cmd = f'kaggle datasets create -p "{SAVE_DIR}"'`: The command to attempt to create a new dataset. The `-p` flag specifies the local folder to upload.

- `subprocess.run(...)` : Executes the command.
- The output is checked for the message "already exists":
 - If it exists, the command `!kaggle datasets version -p "{SAVE_DIR}" -m "..."` is run to upload the data as a new version of the existing dataset. This is the correct way to update an existing Kaggle dataset.
 - If it does not exist, the dataset is created for the first time.
- Finally, the live link and instructions for loading the new dataset using the datasets library are printed.

```
In [ ]: SAVE_DIR = "flat_UI_extended_gemini"

# Run Kaggle dataset creation quietly
result = subprocess.run(
    f'kaggle datasets create -p "{SAVE_DIR}"',
    shell=True,
    capture_output=True,
    text=True
)

# Clear the verbose output from the cell
clear_output(wait=True)

# Check if dataset already exists
if "already exists" in result.stdout.lower():
    print("Dataset exists → creating NEW VERSION")
    # Create a new version quietly
    result_version = subprocess.run(
        f'kaggle datasets version -p "{SAVE_DIR}" -m "10 fresh Gemini captions + original flat UI"',
        shell=True,
        stdout=subprocess.DEVNULL, # hide standard output
        stderr=subprocess.DEVNULL # hide errors (optional)
    )
else:
    print("Dataset created successfully!")

# Print only final info and Link
print("\n Upload done.")
```

```
print(" Dataset ready on Kaggle.")
print("Load later with: kaggle://bhomik7/flat-ui-dataset-extended")
```

Dataset created successfully!

Upload done.

Dataset ready on Kaggle.

Load later with: kaggle://bhomik7/flat-ui-dataset-extended

11. Preview Dataset Images

```
In [ ]: # 11. Get the 20th image (index 19)
image_bytes = full_dataset[19]["image"]
caption = full_dataset[19]["text"]

# Convert to displayable format
img = Image.open(io.BytesIO(image_bytes))

# Display image and caption
display(img)
print(caption)
```

